

student

April 12, 2024

1 Phase 1 Project Submission

- Student name: Connor Anastasio
- Student pace: self paced (part time)
- Scheduled project review date/time: TBD
- Instructor name: Morgan Jones
- Blog post URL: <https://dev.to/connoranastasio/a-history-of-computing-what-led-up-to-ai-4ond>



2 Aircraft Analysis: Low-risk Aviation Investments

2.1 Overview

This project focuses on finding low-risk Aviation aircraft. Investment companies can use this information to determine which aircraft makes and models would be the safest (ie, most profitable) to invest in.

The main approach here will be cleaning our dataset to remove irrelevant information, then manipulate it to uncover relational trends that we can use to gain insight into what factors contribute to accidents.

2.2 Business Problem

Descriptive analysis of this dataset shows that certain makes, models, and types of aircraft appear to be significantly less likely to be involved in an accident than others. Choosing aircraft with one or two engines or those from overall safer companies such as Cessna, Piper, and Beech will allow investors to face significantly reduced risk.

2.3 Data Understanding

We will be using the publicly available [Aviation Accident Data](#) from the United States [National Transport Safety Board](#). It contains very detailed information about each occurrence, including flight information, weather, purpose of flight, and time of day. We will be focusing on eliminating the ‘human element’ and less predictable information from the dataset wherever possible. This approach will enable investors to make informed decisions about what they are investing in. We cannot predict the weather next month or if a flight will need to be made at night or in low visibility, so this information will be less helpful when considering long-term profitability.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
[2]: #load our dataset
#Columns 6, 7, and 28 have mixed data types so we will treat them as strings
columns_to_object = [6, 7, 28]
df = pd.read_csv('data/Aviation_Data.csv', dtype={col: 'str' for col in
    ↪columns_to_object})

#set max columns to none to see information on each column
pd.set_option('display.max_columns', None)
df
```

```
[2]:
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	\
0	20001218X45444	Accident	SEA87LA080	1948-10-24	
1	20001218X45447	Accident	LAX94LA336	1962-07-19	
2	20061025X01555	Accident	NYC07LA005	1974-08-30	
3	20001218X45448	Accident	LAX96LA321	1977-06-19	
4	20041105X01764	Accident	CHI79FA064	1979-08-02	
...	
90343	20221227106491	Accident	ERA23LA093	2022-12-26	
90344	20221227106494	Accident	ERA23LA095	2022-12-26	
90345	20221227106497	Accident	WPR23LA075	2022-12-26	
90346	20221227106498	Accident	WPR23LA076	2022-12-26	
90347	20221230106513	Accident	ERA23LA097	2022-12-29	

	Location	Country	Latitude	Longitude	Airport.Code	\
0	MOOSE CREEK, ID	United States	NaN	NaN	NaN	
1	BRIDGEPORT, CA	United States	NaN	NaN	NaN	
2	Saltville, VA	United States	36.922223	-81.878056	NaN	
3	EUREKA, CA	United States	NaN	NaN	NaN	
4	Canton, OH	United States	NaN	NaN	NaN	
...	
90343	Annapolis, MD	United States	NaN	NaN	NaN	
90344	Hampton, NH	United States	NaN	NaN	NaN	
90345	Payson, AZ	United States	341525N	1112021W	PAN	

90346	Morgan, UT	United States	NaN	NaN	NaN
90347	Athens, GA	United States	NaN	NaN	NaN

	Airport.Name	Injury.Severity	Aircraft.damage	Aircraft.Category	\
0	NaN	Fatal(2)	Destroyed	NaN	
1	NaN	Fatal(4)	Destroyed	NaN	
2	NaN	Fatal(3)	Destroyed	NaN	
3	NaN	Fatal(2)	Destroyed	NaN	
4	NaN	Fatal(1)	Destroyed	NaN	
...	
90343	NaN	Minor	NaN	NaN	
90344	NaN	NaN	NaN	NaN	
90345	PAYSON	Non-Fatal	Substantial	Airplane	
90346	NaN	NaN	NaN	NaN	
90347	NaN	Minor	NaN	NaN	

	Registration.Number	Make	Model	\
0	NC6404	Stinson	108-3	
1	N5069P	Piper	PA24-180	
2	N5142R	Cessna	172M	
3	N1168J	Rockwell	112	
4	N15NY	Cessna	501	
...	
90343	N1867H	PIPER	PA-28-151	
90344	N2895Z	BELLANCA	7ECA	
90345	N749PJ	AMERICAN CHAMPION AIRCRAFT	8GCBC	
90346	N210CU	CESSNA	210N	
90347	N9026P	PIPER	PA-24-260	

	Amateur.Built	Number.of.Engines	Engine.Type	FAR.Description	\
0	No	1.0	Reciprocating	NaN	
1	No	1.0	Reciprocating	NaN	
2	No	1.0	Reciprocating	NaN	
3	No	1.0	Reciprocating	NaN	
4	No	NaN	NaN	NaN	
...	
90343	No	NaN	NaN	091	
90344	No	NaN	NaN	NaN	
90345	No	1.0	NaN	091	
90346	No	NaN	NaN	091	
90347	No	NaN	NaN	091	

	Schedule	Purpose.of.flight	Air.carrier	Total.Fatal.Injuries	\
0	NaN	Personal	NaN	2.0	
1	NaN	Personal	NaN	4.0	
2	NaN	Personal	NaN	3.0	
3	NaN	Personal	NaN	2.0	

4	NaN	Personal	NaN	1.0
...
90343	NaN	Personal	NaN	0.0
90344	NaN	NaN	NaN	0.0
90345	NaN	Personal	NaN	0.0
90346	NaN	Personal	MC CESSNA 210N LLC	0.0
90347	NaN	Personal	NaN	0.0

	Total.Serious.Injuries	Total.Minor.Injuries	Total.Uninjured	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	NaN	NaN	NaN	
3	0.0	0.0	0.0	
4	2.0	NaN	0.0	
...	
90343	1.0	0.0	0.0	
90344	0.0	0.0	0.0	
90345	0.0	0.0	1.0	
90346	0.0	0.0	0.0	
90347	1.0	0.0	1.0	

	Weather.Condition	Broad.phase.of.flight	Report.Status	Publication.Date
0	UNK	Cruise	Probable Cause	NaN
1	UNK	Unknown	Probable Cause	19-09-1996
2	IMC	Cruise	Probable Cause	26-02-2007
3	IMC	Cruise	Probable Cause	12-09-2000
4	VMC	Approach	Probable Cause	16-04-1980
...
90343	NaN	NaN	NaN	29-12-2022
90344	NaN	NaN	NaN	NaN
90345	VMC	NaN	NaN	27-12-2022
90346	NaN	NaN	NaN	NaN
90347	NaN	NaN	NaN	30-12-2022

[90348 rows x 31 columns]

```
[3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 90348 entries, 0 to 90347
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Event.Id	88889 non-null	object
1	Investigation.Type	90348 non-null	object
2	Accident.Number	88889 non-null	object
3	Event.Date	88889 non-null	object
4	Location	88837 non-null	object

```

5   Country                88663 non-null object
6   Latitude               34382 non-null object
7   Longitude              34373 non-null object
8   Airport.Code           50132 non-null object
9   Airport.Name           52704 non-null object
10  Injury.Severity         87889 non-null object
11  Aircraft.damage         85695 non-null object
12  Aircraft.Category       32287 non-null object
13  Registration.Number     87507 non-null object
14  Make                   88826 non-null object
15  Model                   88797 non-null object
16  Amateur.Built           88787 non-null object
17  Number.of.Engines       82805 non-null float64
18  Engine.Type             81793 non-null object
19  FAR.Description         32023 non-null object
20  Schedule                12582 non-null object
21  Purpose.of.flight       82697 non-null object
22  Air.carrier             16648 non-null object
23  Total.Fatal.Injuries    77488 non-null float64
24  Total.Serious.Injuries  76379 non-null float64
25  Total.Minor.Injuries    76956 non-null float64
26  Total.Uninjured         82977 non-null float64
27  Weather.Condition       84397 non-null object
28  Broad.phase.of.flight   61724 non-null object
29  Report.Status           82505 non-null object
30  Publication.Date        73659 non-null object
dtypes: float64(5), object(26)
memory usage: 21.4+ MB

```

Our dataset contains an impressive amount of individual entries and information on flights dating back to 1948.

As we are primarily concerned with the objective safety of each make and model, we can begin data cleaning and investigation as to which aspects will not be useful for this.

2.4 Data Preparation

2.4.1 Data Cleaning

```

[4]: #normalize column names for ease of use
df.columns = df.columns.str.lower().str.replace('.', '_')
df['make'].replace('Air Tractor', 'Air Tractor Inc', inplace=True)

```

At a glance, certain columns will not be useful for our analysis and we can drop them now. Others will need to be looked at further before we make our judgment.

```

[5]: #drop obvious unnecessary columns
df.drop(columns = ['latitude', 'longitude', 'event_id'], inplace=True)

```

```
[6]: #Examine Broad Phase of Flight
print(df['broad_phase_of_flight'].value_counts())
print('broad phase null counts:', df['broad_phase_of_flight'].isnull().sum())
```

```
broad_phase_of_flight
Landing      15428
Takeoff      12493
Cruise       10269
Maneuvering   8144
Approach      6546
Climb         2034
Taxi          1958
Descent       1887
Go-around    1353
Standing      945
Unknown       548
Other         119
Name: count, dtype: int64
broad phase null counts: 28624
```

```
[7]: df.drop(columns = ['broad_phase_of_flight'], inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 90348 entries, 0 to 90347
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   investigation_type     90348 non-null  object
1   accident_number        88889 non-null  object
2   event_date             88889 non-null  object
3   location                88837 non-null  object
4   country                88663 non-null  object
5   airport_code           50132 non-null  object
6   airport_name           52704 non-null  object
7   injury_severity        87889 non-null  object
8   aircraft_damage        85695 non-null  object
9   aircraft_category      32287 non-null  object
10  registration_number     87507 non-null  object
11  make                   88826 non-null  object
12  model                  88797 non-null  object
13  amateur_built          88787 non-null  object
14  number_of_engines       82805 non-null  float64
15  engine_type            81793 non-null  object
16  far_description         32023 non-null  object
17  schedule               12582 non-null  object
18  purpose_of_flight       82697 non-null  object
19  air_carrier            16648 non-null  object
```

```

20 total_fatal_injuries    77488 non-null float64
21 total_serious_injuries  76379 non-null float64
22 total_minor_injuries    76956 non-null float64
23 total_uninjured         82977 non-null float64
24 weather_condition       84397 non-null object
25 report_status           82505 non-null object
26 publication_date        73659 non-null object
dtypes: float64(5), object(22)
memory usage: 18.6+ MB

```

Broad Phase of Flight could be interesting to consider further but unfortunately it is missing far too many values to be useful in our analysis.

Additionally, investing in a non-professionally made aircraft is undoubtedly a risky and unpredictable investment. We should only consider accident data for professionally made aircraft. We will have to also drop the rows where it is unknown as we cannot confirm it was not amateur built:

```

[8]: #Remove All Non-Accident Investigation Type Rows, and all amateur built aircraft
df = df[df['investigation_type'] == 'Accident']
df = df[df['amateur_built'] == 'No']
df.drop(columns = ['investigation_type', 'accident_number', 'airport_code',
    ↳ 'airport_name', 'aircraft_category', 'registration_number',
    ↳ 'amateur_built', 'far_description', 'schedule',
    ↳ 'purpose_of_flight', 'air_carrier', 'publication_date'], inplace=True)

```

2.4.2 Data Engineering

It will be helpful to create a few of our own columns for analysis.

```

[9]: df['total_injuries'] = df['total_fatal_injuries'] +
    ↳ df['total_serious_injuries'] + df['total_minor_injuries']
df['total_non_fatal_injuries'] = df['total_serious_injuries'] +
    ↳ df['total_minor_injuries']
df

```

```

[9]:
   event_date      location      country injury_severity \
0   1948-10-24  MOOSE CREEK, ID  United States      Fatal(2)
1   1962-07-19  BRIDGEPORT, CA  United States      Fatal(4)
2   1974-08-30   Saltville, VA  United States      Fatal(3)
3   1977-06-19    EUREKA, CA   United States      Fatal(2)
4   1979-08-02    Canton, OH   United States      Fatal(1)
...         ...           ...           ...           ...
90343  2022-12-26  Annapolis, MD  United States      Minor
90344  2022-12-26   Hampton, NH  United States      NaN
90345  2022-12-26   Payson, AZ   United States    Non-Fatal
90346  2022-12-26   Morgan, UT   United States      NaN
90347  2022-12-29   Athens, GA   United States      Minor

```

```

aircraft_damage      make      model \

```

0	Destroyed		Stinson	108-3
1	Destroyed		Piper	PA24-180
2	Destroyed		Cessna	172M
3	Destroyed		Rockwell	112
4	Destroyed		Cessna	501
...
90343	NaN		PIPER	PA-28-151
90344	NaN		BELLANCA	7ECA
90345	Substantial	AMERICAN CHAMPION	AIRCRAFT	8GCBC
90346	NaN		CESSNA	210N
90347	NaN		PIPER	PA-24-260

	number_of_engines	engine_type	total_fatal_injuries \
0	1.0	Reciprocating	2.0
1	1.0	Reciprocating	4.0
2	1.0	Reciprocating	3.0
3	1.0	Reciprocating	2.0
4	NaN	NaN	1.0
...
90343	NaN	NaN	0.0
90344	NaN	NaN	0.0
90345	1.0	NaN	0.0
90346	NaN	NaN	0.0
90347	NaN	NaN	0.0

	total_serious_injuries	total_minor_injuries	total_uninjured \
0	0.0	0.0	0.0
1	0.0	0.0	0.0
2	NaN	NaN	NaN
3	0.0	0.0	0.0
4	2.0	NaN	0.0
...
90343	1.0	0.0	0.0
90344	0.0	0.0	0.0
90345	0.0	0.0	1.0
90346	0.0	0.0	0.0
90347	1.0	0.0	1.0

	weather_condition	report_status	total_injuries \
0	UNK	Probable Cause	2.0
1	UNK	Probable Cause	4.0
2	IMC	Probable Cause	NaN
3	IMC	Probable Cause	2.0
4	VMC	Probable Cause	NaN
...
90343	NaN	NaN	1.0
90344	NaN	NaN	0.0

90345	VMC	NaN	0.0
90346	NaN	NaN	0.0
90347	NaN	NaN	1.0

	total_non_fatal_injuries
0	0.0
1	0.0
2	NaN
3	0.0
4	NaN
...	...
90343	1.0
90344	0.0
90345	0.0
90346	0.0
90347	1.0

[76520 rows x 17 columns]

We need to ensure the date column is properly formatted for all 76,520 remaining rows. We can do this by writing a simple for loop to ensure all of the values can be converted to standard Pandas datetime. If this code does not return an error, we won't have to worry.

```
[10]: def check_date_format(date):
    try:
        pd.to_datetime(date)
        return True
    except ValueError:
        return False

    # Check if all dates in the column are properly formatted
    all_dates_valid = df['event_date'].map(check_date_format).all()

    if all_dates_valid:
        print("All dates are properly formatted.")
    else:
        print("Some dates have formatting issues.")
```

All dates are properly formatted.

```
[11]: #create a new column of make + model
df['make_model'] = df['make'] + '_' + df['model']

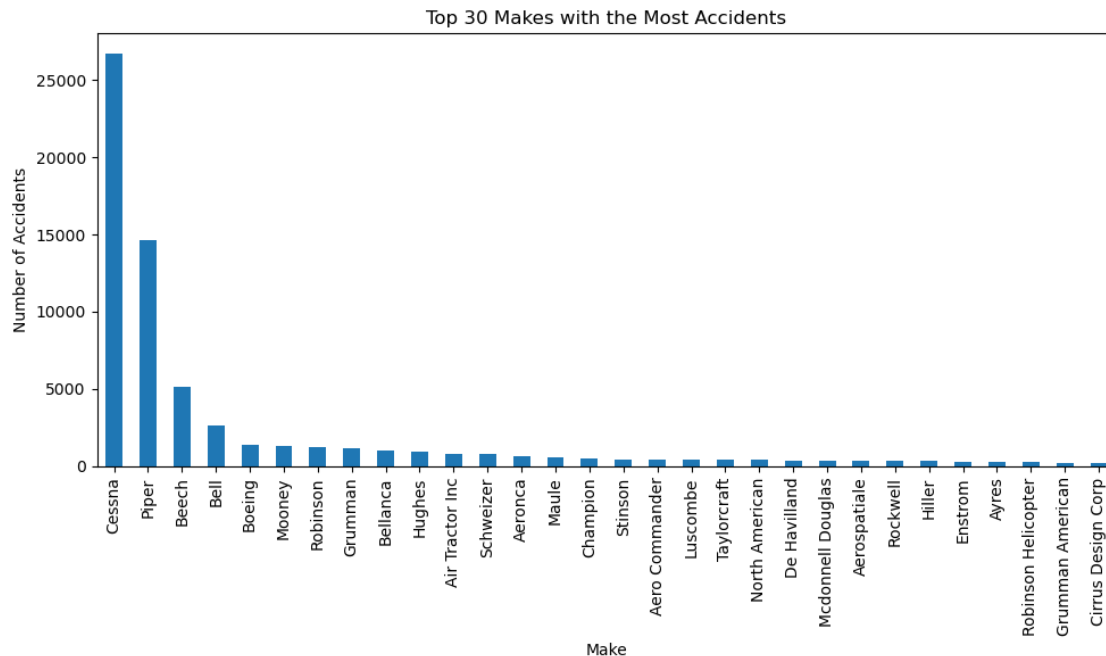
#Format 'make' column and describe its properties
df['make'] = df['make'].str.title()
df['make'].describe()
```

```
[11]: count      76500
      unique      1880
      top        Cessna
      freq       26695
      Name: make, dtype: object
```

As we can see, there are 1880 different Makes, and many appear only once. This isn't very helpful in looking for data trends, so let's create a histogram with the top 30 most frequently occurring makes in our dataset:

```
[12]: #Create a variable to store the top 30 most frequently occurring makes:
      top_30_makes = df['make'].value_counts().nlargest(30)

      # Plotting the histogram for top 30 makes
      plt.figure(figsize=(10, 6))
      top_30_makes.plot(kind='bar')
      plt.title('Top 30 Makes with the Most Accidents')
      plt.xlabel('Make')
      plt.ylabel('Number of Accidents')
      plt.xticks(rotation=90)
      plt.tight_layout()
      plt.show()
```



Cessna has the most accidents by a wide margin, but we will need additional and deeper analysis in order to make assessments on this.

```
[13]: #analyze frequency
df['year'] = pd.to_datetime(df['event_date']).dt.year
accidents_since_2000 = df[df['year'] >= 2000]
accidents_since_2000
```

```
[13]:
```

	event_date	location	country	injury_severity	\
47675	2000-01-01	HOMESTEAD, FL	United States	Non-Fatal	
47676	2000-01-01	MONTEAGLE, TN	United States	Fatal(2)	
47677	2000-01-02	VICTORVILLE, CA	United States	Non-Fatal	
47678	2000-01-02	DOS PALOS, CA	United States	Non-Fatal	
47679	2000-01-02	CORNING, AR	United States	Non-Fatal	
...	
90343	2022-12-26	Annapolis, MD	United States	Minor	
90344	2022-12-26	Hampton, NH	United States	NaN	
90345	2022-12-26	Payson, AZ	United States	Non-Fatal	
90346	2022-12-26	Morgan, UT	United States	NaN	
90347	2022-12-29	Athens, GA	United States	Minor	

	aircraft_damage	make	model	\
47675	Substantial	Cessna	550	
47676	Destroyed	Bellanca	BL-17-30A	
47677	Substantial	Cessna	172G	
47678	Substantial	Cessna	172A	
47679	Substantial	Piper	PA-46-310P	
...	
90343	NaN	Piper	PA-28-151	
90344	NaN	Bellanca	7ECA	
90345	Substantial	American Champion Aircraft	8GCBC	
90346	NaN	Cessna	210N	
90347	NaN	Piper	PA-24-260	

	number_of_engines	engine_type	total_fatal_injuries	\
47675	2.0	Turbo Fan	0.0	
47676	1.0	Reciprocating	2.0	
47677	1.0	Reciprocating	0.0	
47678	1.0	Reciprocating	0.0	
47679	1.0	Turbo Prop	0.0	
...	
90343	NaN	NaN	0.0	
90344	NaN	NaN	0.0	
90345	1.0	NaN	0.0	
90346	NaN	NaN	0.0	
90347	NaN	NaN	0.0	

	total_serious_injuries	total_minor_injuries	total_uninjured	\
47675	0.0	0.0	3.0	
47676	0.0	0.0	0.0	

47677	0.0	0.0	2.0
47678	0.0	1.0	0.0
47679	0.0	0.0	5.0
...
90343	1.0	0.0	0.0
90344	0.0	0.0	0.0
90345	0.0	0.0	1.0
90346	0.0	0.0	0.0
90347	1.0	0.0	1.0

	weather_condition	report_status	total_injuries \
47675	VMC	Probable Cause	0.0
47676	IMC	Probable Cause	2.0
47677	VMC	Probable Cause	0.0
47678	VMC	Probable Cause	1.0
47679	VMC	Probable Cause	0.0
...
90343	NaN	NaN	1.0
90344	NaN	NaN	0.0
90345	VMC	NaN	0.0
90346	NaN	NaN	0.0
90347	NaN	NaN	1.0

	total_non_fatal_injuries	make_model	year
47675	0.0	Cessna_550	2000
47676	0.0	Bellanca_BL-17-30A	2000
47677	0.0	Cessna_172G	2000
47678	1.0	Cessna_172A	2000
47679	0.0	Piper_PA-46-310P	2000
...
90343	1.0	PIPER_PA-28-151	2022
90344	0.0	BELLANCA_7ECA	2022
90345	0.0	AMERICAN CHAMPION AIRCRAFT_8GCBC	2022
90346	0.0	CESSNA_210N	2022
90347	1.0	PIPER_PA-24-260	2022

[34210 rows x 19 columns]

We will focus on the 30 most common aircraft makes for our analysis; these are the companies with enough instances to be statistically relevant.

While older planes are still in service, flight safety and security has improved drastically in our post-9/11 world. For this reason it is worth looking into whether there has been a noticeable decrease in accidents since:

```
[14]: top_30_aircraft = df['make'].value_counts().nlargest(30).index.tolist()

# Filtering DataFrame for 30 most prevalent aircraft makes
```

```

df_top_30 = df[df['make'].isin(top_30_aircraft)]


# Grouping data by make and year, and counting occurrences
df_top_30_since_2000 = df[(df['make'].isin(top_30_aircraft)) & (df['year'] >=
↳ 2000)]
grouped = df_top_30_since_2000.groupby(['make', 'year']).size().
↳ reset_index(name='accidents')

# Creating a line chart for each aircraft make
plt.figure(figsize=(12, 9))

for make in top_30_aircraft:
    make_data = grouped[grouped['make'] == make]
    plt.plot(make_data['year'], make_data['accidents'], label=make)

plt.title('Accidents by Year for Top 30 Aircraft Makes')
plt.xlabel('Year')
plt.ylabel('Number of Accidents')
plt.legend(bbox_to_anchor=(1.05, 1), loc='upper left', borderaxespad=0)
plt.show()

```



While this provides us with interesting information, it is not the full picture of what is happening. The dataset does not include the amount of planes in service for each company by year, so we don't know the percentage of each company's models that are involved in accidents. Still, it is noteworthy that every aircraft maker is on a downward trend in number of accidents. This implies it is a safer

and potentially more profitable to invest in essentially any aircraft company than it was even 15 years ago.

Although we do not have the ability to create an “in-use” aircraft ratio by company, we thankfully have plenty of other useful information in our dataset, and approaches we can take.

The relationship between the number of engines and the total injuries will be instrumental in our analysis. To analyze this, we will create two bar plots:

```
[15]: # Calculate frequency of instances for each number of engines
instance_counts = df['number_of_engines'].value_counts().sort_index()

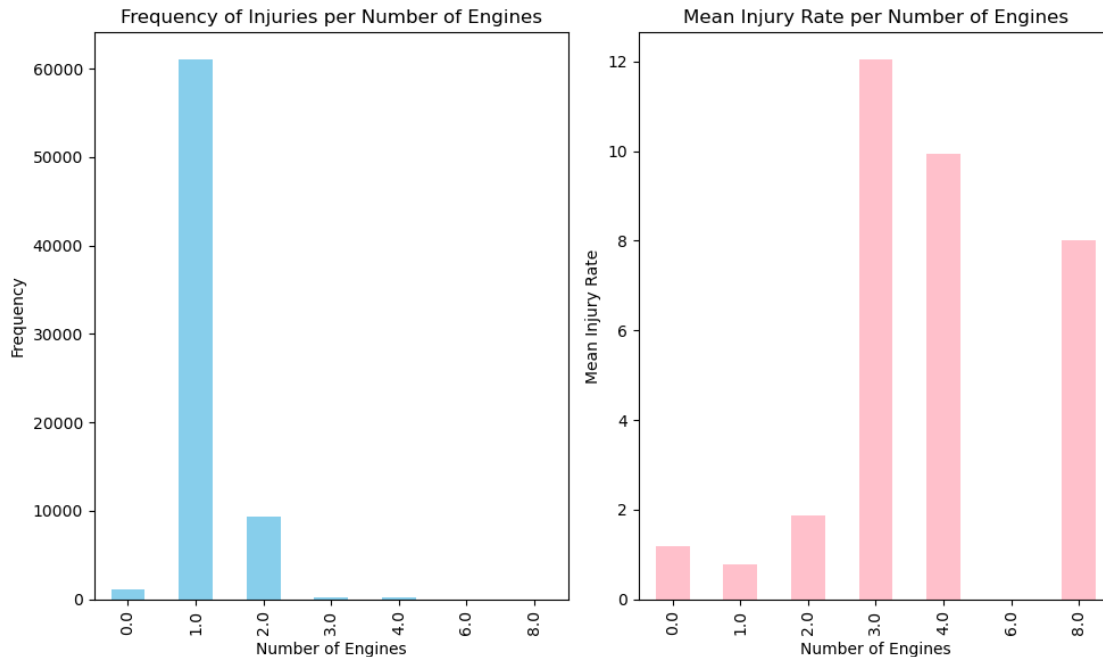
# Calculate mean injury rate for each number of engines
injury_rates = df.groupby('number_of_engines')['total_injuries'].mean().
    ↪sort_index()

plt.figure(figsize=(10, 6))

# Plot the bar plots side-by-side
plt.subplot(1, 2, 1)
instance_counts.plot(kind='bar', color='skyblue')
plt.title('Frequency of Injuries per Number of Engines')
plt.xlabel('Number of Engines')
plt.ylabel('Frequency')

plt.subplot(1, 2, 2)
injury_rates.plot(kind='bar', color='pink')
plt.title('Mean Injury Rate per Number of Engines')
plt.xlabel('Number of Engines')
plt.ylabel('Mean Injury Rate')

plt.tight_layout()
plt.show()
```



The left subplot displays a bar plot showing the frequency of instances for each number of engines, and the right shows a bar plot representing the mean injury rates for the corresponding number of engines.

Using these plots, it is easy to see that injury rates tend to increase as the number of engines does. In other words, aircraft with 1 or 2 engines appear to be the safest.


Notably, single engine aircraft have the second lowest mean injury rate while making up nearly all of the instances in our dataset. We will focus our analysis on these, and determine the safest make/model with one engine:

```
[16]: # Filtering df_top_30 for entries where the number of engines is one
df_top_30_single_engine = df_top_30[df_top_30['number_of_engines'] == 1]

# Grouping data by 'make' for single-engine aircraft (top 30 makes) and
# calculating mean injury rates
make_injury_rates_single_engine = df_top_30_single_engine.
    .groupby('make')['total_injuries'].mean().sort_values()

# Plotting bar plot
make_injury_rates_single_engine.plot(kind='bar', figsize=(12, 6),
    color='skyblue')
plt.title('Mean Injury Rates by Make (Single-Engine Aircraft - Top 30 Makes)')
plt.xlabel('Make')
plt.ylabel('Mean Injury Rate')
plt.xticks(rotation=90)
```

```
plt.show()
```



There are 22 makes with an injury rate below 1.0. This is excellent, as further restricting to these will still give us plenty of options to invest in.

In order to see if there is a single company with the overall best safety rating, we will examine the models in a similar fashion (using the top 30 most common again). In order to do this, we will find each model with an average injury rate less than 1, add each by make, and then plot them:

```
[17]: # Create a variable with the top 30 most common makes
top_30_makes = df['make'].value_counts().nlargest(30).index.tolist()

# Filter the DataFrame for the top 30 makes
df_top_30_makes = df[df['make'].isin(top_30_makes)]

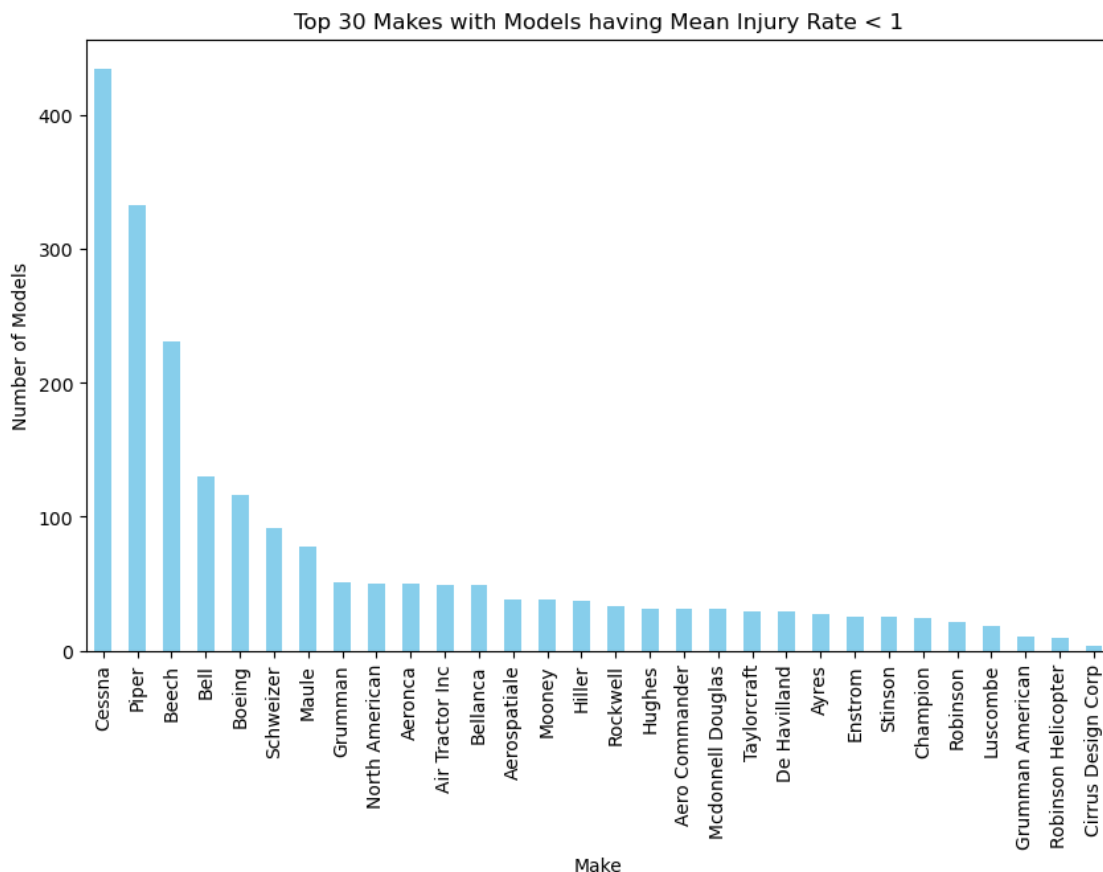
# Calculate the mean injury rate per model for the top 30 makes
mean_injury_rate_per_model_top_30 = df_top_30_makes.groupby(['make', '
    ↳model'])['total_injuries'].mean()

# Filter models with a mean injury rate < 1 for the top 30 makes
models_with_mean_injury_lt_1_top_30 =
    ↳mean_injury_rate_per_model_top_30[mean_injury_rate_per_model_top_30 < 1].
    ↳reset_index()
```



```
# Count the number of models with mean injury rate < 1 for each make among the
↳ top 30 makes
models_lt_1_count_per_make_top_30 = models_with_mean_injury_lt_1_top_30.
↳ groupby('make')['model'].count().sort_values(ascending=False)

# Plotting the top 30 makes with the count of models having mean injury rate < 1
plt.figure(figsize=(10, 6))
models_lt_1_count_per_make_top_30.plot(kind='bar', color='skyblue')
plt.title('Top 30 Makes with Models having Mean Injury Rate < 1')
plt.xlabel('Make')
plt.ylabel('Number of Models')
plt.xticks(rotation=90)
plt.show()
```



Cessna has the most models with the the lowest injury rates. Hooray!

2.4.3 Weather Influence

It will be useful to see if there is a noticeable relationship between poor visibility and accident statistics.

```

[18]: # Filter data for 'IMC' and 'VMC' while dropping null values in 'total_injuries'
imc_data = df[(df['weather_condition'] == 'IMC') & (df['total_injuries'].
    ↪notnull())].copy()
vmc_data = df[(df['weather_condition'] == 'VMC') & (df['total_injuries'].
    ↪notnull())].copy()

# Group counts of 4.0 injuries or more as '4+'
imc_data.loc[:, 'total_injuries_grouped'] = imc_data['total_injuries'].
    ↪apply(lambda x: '4.0+' if x >= 4.0 else str(x))
vmc_data.loc[:, 'total_injuries_grouped'] = vmc_data['total_injuries'].
    ↪apply(lambda x: '4.0+' if x >= 4.0 else str(x))

# Calculate value counts for 'IMC' and 'VMC' with grouping
imc_counts = imc_data['total_injuries_grouped'].value_counts()
vmc_counts = vmc_data['total_injuries_grouped'].value_counts()

# Convert counts to dictionary for pie chart plotting
imc_counts_dict = imc_counts.to_dict()
vmc_counts_dict = vmc_counts.to_dict()

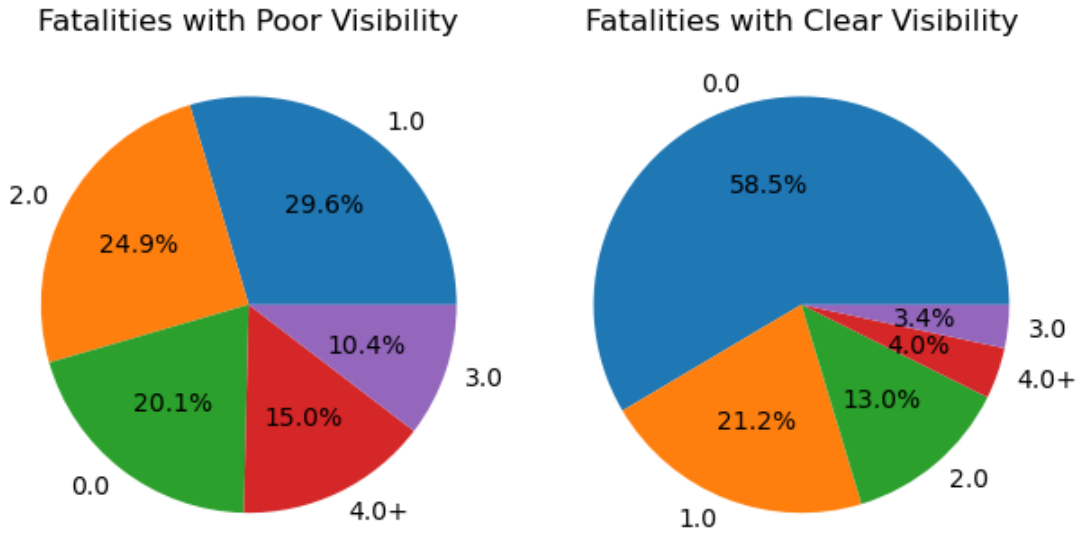
# Replace count of '4+' if it doesn't exist in 'IMC' or 'VMC' counts
if '4.0+' not in imc_counts_dict:
    imc_counts_dict['4.0+') = 0
if '4.0+' not in vmc_counts_dict:
    vmc_counts_dict['4.0+') = 0

# Plot pie chart for 'IMC'
plt.subplot(1, 2, 1)
plt.pie(imc_counts_dict.values(), labels=imc_counts_dict.keys(), autopct='%1.
    ↪1f%%')
plt.title('Fatalities with Poor Visibility')

# Plot pie chart for 'VMC'
plt.subplot(1, 2, 2)
plt.pie(vmc_counts_dict.values(), labels=vmc_counts_dict.keys(), autopct='%1.
    ↪1f%%')
plt.title('Fatalities with Clear Visibility')

plt.tight_layout()
plt.show()

```



It is apparent from these pie charts that aircraft flying in poor visibility conditions (requiring pilots to rely solely on instruments) leads to a significant increase in casualties. It is highly advised to avoid investing in planes that fly routes that have poor visibility.

2.5 Conclusion:

- In terms of overall safety ratings for investing, single engine planes should make up the majority of investments.
- Investing in a single well-known company may be a safe approach to start with. Cessna is a great choice, as our dataset has provided ample information for model choices.
- Avoiding routes with poor visibility and providing further training pilots to deal with adverse conditions would likely aid in reducing accidents and increase profit.

2.6 Future:

- Compare cost of the safest aircraft models to determine which can bring the largest return on investment
- Gathering data on Aircraft sales and in-use data for individual aircraft would allow us to make more informative individual model recommendations